

Sheet3

07.Jan.2020 Abraham Gutierrez && Huan Chen

Exercise 1, inner product

(1) run the code and get the following result:

```
g++ main.o mylib.o -g -flto -fopenmp -o main.GCC_
./main.GCC_
Number of available processors: 64
The return of command 'omp_in_parallel': 1
Checking command line parameters for: -n <number>

N = 40000000

#####
Code : ./main.GCC_
Compiler: Gnu 9.2.0 C++ standard: 201703
Parallel: OpenMP 4.5 ---> 64 Threads
Date : Dec 15 2019 20:37:52
#####
C++: Hello World from thread 0 / 64
C++: Hello World from thread 38 / 64
C++: Hello World from thread 25 / 64
C++: Hello World from thread 18 / 64
C++: Hello World from thread 20 / 64
C++: Hello World from thread 14 / 64
C++: Hello World from thread 61 / 64
C++: Hello World from thread 23 / 64
C++: Hello World from thread 8 / 64
C++: Hello World from thread 62 / 64
C++: Hello World from thread 52 / 64
C++: Hello World from thread 26 / 64
C++: Hello World from thread 58 / 64
.
.
.
C++: Hello World from thread 59 / 64

64 threads have been started.
Memory allocation
0.6 GByte Memory allocated

Start Benchmarking

<x,y> = 4e+07
```

Total time in sec. :2.6
Time/loop in sec. : 0.013
GFLOPS : 5.7
GiByte/s : 46

Try the reduction with an STL-vektor
done

2016 2080 2144 2208 2272 2336 2400 2464 2528 2592 2656 2720 2784 2848 2912 2976
3040 3104 3168 3232 3296 3360 3424 3488 3552 3616 3680 3744 3808 3872 3936 4000
4064 4128 4192 4256 4320 4384 4448 4512 4576 4640 4704 4768 4832 4896 4960 5024
5088 5152 5216 5280 5344 5408 5472 5536 5600 5664 5728 5792 5856 5920 5984 6048
6112 6176 6240 6304 6368 6432 6496 6560 6624 6688 6752 6816 6880 6944 7008 7072
7136 7200 7264 7328 7392 7456 7520 7584 7648 7712 7776 7840 7904 7968 8032 8096
8160 8224 8288 8352

(2) Try several schedule types and chunk sizes: **schedule(kind , chunk size)**

$N = 4 \cdot 10^7$, NLOOPS = 200.

	Total time (sec)	GFLOPS	GiByte/s
No schedule	2.6	5.7	46
schedule(auto)	2.5	6	48
schedule(static)	2.4	6.2	50
schedule(static, 250)	4.8	3.1	25
schedule(static, 25000)	2.7	5.6	45
schedule(dynamic)	/	/	/
schedule(dynamic, 250)	4.9	3	24
schedule(dynamic, 25000)	2.6	5.8	46
schedule(guided, 250)	2.6	5.7	46
schedule(guided, 25000)	2.6	5.8	46

Conclusion:

- It's important to find the proper chunk size, e.g. 25000 is better than 250
- For dynamic, if we do not specify chunk size, it sets to defaults (one), which is seriously inefficient in this case.
- In this simple for loop, auto or default schedule is behaving well enough.

(3) Calculate the speedup for different number of cores: **omp set num threads()**

$N = 4 \times 10^7$, NLOOPS = 200;

No schedule.

The computer used to do the testing has 32 cores with 2 threads for each core i.e. 64 threads.

Numbet of threads	Timing in sec	GFLOPS	GiByte/s
1	5.534	2.693	21.54
2	4.024	3.703	29.62
4	2.721	5.476	43.8
8	3.014	4.943	39.55
16	3.01	4.951	39.6
32	2.852	5.225	41.8
64	2.946	5.057	40.46

```
cout << "Number of available processors: " << omp_get_num_procs() << endl;  
Number of available processors: 64
```

The return of command 'omp_in_parallel': 1

The result means: number of nested active parallel regions (active-levels-var) is larger than zero (otherwise returns 0).

Appending without order VS appending with order

Number of Loops: 200

n = 3000000

Number of Threads	Withour order, total time (s)	With order, total time (s)
1	0.6472	0.6385
2	2.568	2.074
4	8.423	10.26
8	25.41	27.22

Exercise 2, data I/O

First all, we produced a file.txt with $2 \cdot 10^8$ double elements inside;

We didn't parallize the data 'reading' process (does that worth it and how), time for that is: 37.8032s

So the following comparison is only for the (max, min, mean values) calculation part, NLOOPS = 20, correct result is: 1 1000 500.5 369.5 133.6 288.7

	Total time in sec	Time per loop in sec
Non-parallized	256.4	12.82
Parallized	8.025	0.4013

Efficiency = $256.4 / 8.025 / 64 = 0.4992$.

Exercise 3, count_goldbach

Version 1: g++ -O3, number of used threads: nThreads = 64

	Non-parallelization time in sec, t1	Parallelization time in sec, t2	Efficiency= $t1/t2/nThreads$
n=10,000	0.001289	0.009314	0.0022
n=100,000	0.164	0.02632	0.097
n=400,000	1.361	0.1348	0.16
n=500,000	2.042	0.1808	0.18
n=1,000,000	7.792	0.6222	0.20
n=2,000,000	28.25	3.246	0.14
n=5,000,000	155.6	28.73	0.085

Version 2: g++ -O3

n=10.000.000	-O3
1	555.5 sec
2	273.9 sec
4	136 sec
8	69.51 sec

16	42.7 sec
32	24.29 sec
64	24.33 sec

Exercise 4

Matrix-Vector Multiplication

M times N matrix: M = 60000, N = 9000

Number of Loops: 1.5e+02

Matrix-Matrix Multiplication

Matrix Dimensions (M times L and L times N): M = 2000, L = 2000, N = 2000

Number of Loops: 5

Polynomial Evaluation

Size of vector x: 5000

Polynomial size (N) : 200000

Number of Loops: 10

TimesMatrixVector = [91 48 28 18 18 22 20];

TimesMatrixMatrix = [19 15 9.8 7.3 4.3 2.2 1.6];

TimesPolynomialEvaluation = [14 7.2 3.6 2 1.1 0.61 0.54];

EfficiencyMatrixVector = [1 0.96 0.81 0.65 0.32 0.13 0.07];

EfficiencyMatrixMatrix = [1 0.64 0.49 0.33 0.28 0.27 0.19];

EfficiencyPolynomialEvaluation = [1 0.95 0.95 0.88 0.77 0.7 0.4];

GFlopsMatrixVector = [0.011 0.021 0.036 0.057 0.056 0.046 0.05];

GFlopsMatrixMatrix = [0.77 0.98 1.5 2.1 3.5 6.7 9.4];

GFlopsPolynomialEvaluation = [0.2 0.39 0.77 1.4 2.5 4.6 5.2];

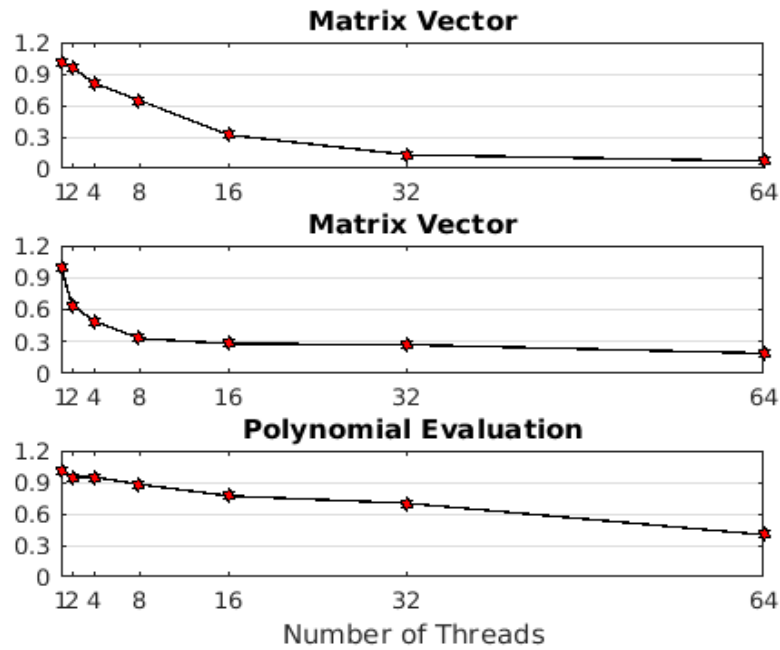
GipPerSecMatrixVector = [0.18 0.34 0.57 0.92 0.9 0.74 0.8];

GipPerSecMatrixMatrix = [12 16 24 33 56 1.1e+02 1.5e+02];

GipPerSecPolynomialEvaluation = [3.8 7.2 14 27 47 86 97];

NumberofThreads = [1 2 4 8 16 32 64];

Efficiency plot



Two versions of Polynomial Evaluation comparison

Size of vector x: 5000

Polynomial size (N) : 200000

Number of Loops: 10

Number of Threads	Parallized outer loop, total time (s)	Parallized inner loop, total time (s)
1	13.4629	15.0908
2	6.79469	7.69313
4	3.47619	4.03413
8	1.97234	2.30301
16	1.20596	1.50045
32	0.672577	1.44844
64	0.514111	2.19934

Exercise 5, Jacobi

Non-parallelized version

Make run, get the following result:

Intervalls: 100 x 100

Start Jacobi solver for 10201 d.o.f.s

aver. Jacobi rate : 0.997922 (1000 iter)

final error: 0.124971 (rel) 0.000194029 (abs)

JacobiSolve: timing in sec. : 0.080958

ASCII file square_100.txt opened

17361 2 34320 3

Start Jacobi solver for 17361 d.o.f.s

aver. Jacobi rate : 0.998401 (1000 iter)

final error: 0.201744 (rel) 0.000265133 (abs)

JacobiSolve: timing in sec. : 0.189638

Parallelized version

Intervalls: 100 x 100, -O3, 1000 iter

Number of Threads	Time1, s	Time2, s
1	0.0774207	0.189079
2	0.0867711	0.1952
4	0.0979033	0.223702
8	0.0935012	0.223959
16	0.12996	0.234558
32	0.151673	0.22109
64	0.264242	0.303877

No improvement, perhaps calculation amount is too small.

Intervalls: 1000 x 1000, -O3, 1000 iter

Number of Threads	Time1, s	Time2, s
1	9.65759	0.188669
2	6.32916	0.175074
4	4.51213	0.120928
8	4.42104	0.0979477
16	3.79665	0.0879493
32	3.63878	0.0947926

64	3.76486	0.203179
----	---------	----------

When threads number are 2 and 4, the efficiencies are best. When the processes number increase, the communication time between caches of processors increase for sharing data.